




*This work: all 5 runs share LoRA rank 128, so differences here are backbone architecture only.
gpt_quantum reuses the LFM2.5-350M backbone and LoRA config unchanged, so its bar coincides with it.
The paper does not publish its LoRA rank / target modules -- the two groups are not on a common configuration. Read them as two separate groups; do not compute cross-group deltas.*