




*This work: peak memory over one real training step, measured on Apple MPS (reports/measure\_vram\_usage.py) -- NOT measured on the RTX 4090 that ran the 100-epoch training, where no VRAM profiling was instrumented.  
Paper: Fig. 8(b), 2x NVIDIA RTX 6000. Different hardware on both sides; do not compute cross-group deltas.*